

# PyTorch practice

Multilayer Perceptron (MLP)

---

Matteo Boschini, Lorenzo Bonicelli

November 25, 2022

University of Modena and Reggio Emilia

**Multilayer perceptron (MLP)**

**References**

# Multilayer perceptron (MLP)

---

For the purpose of this lecture, we'll stick to the task of image classification.  
Let's assume we have a training dataset of  $N$  images

$$x_i \in \mathbb{R}^D, i = 1, \dots, N$$


that we want to classify into  $K$  distinct classes.

Thus, training set is made by couples:

$$(x_i, y_i), \text{ where } y_i \in \{1, \dots, K\}$$

Our goal is to define a function  $f : \mathbb{R}^D \mapsto \mathbb{R}^K$  that maps images to class scores.

When we connect an ensemble of neurons in a graph is when the magic happens and we get an actual **Multilayer perceptron (MLP)**.



Neural networks are arranged in **layers**, with one *input layer*, one *output layer* and  $N$  *hidden layers* in the middle.


*The network depicted here has a total of 47 learnable parameters. Does this make sense to you?*

The 4-layer network previously depicted can be simply expressed as:

$$out = \phi(\mathbf{W}_3\phi(\mathbf{W}_2\phi(\mathbf{W}_1\mathbf{x}))) \quad (1)$$


where:

- $\phi$  is the activation function
- $\mathbf{x} \in \mathbb{R}^6$  is the input
- $\mathbf{W}_1 \in \mathbb{R}^{4 \times 6}$  are the weights of first layer
- $\mathbf{W}_2 \in \mathbb{R}^{3 \times 4}$  are the weights of second layer
- $\mathbf{W}_3 \in \mathbb{R}^{1 \times 3}$  are the weights of third layer



Notice that to ease the notation biases have been incorporated into weight matrices  $\mathbf{W}$ .

**Activation functions** are non-linear functions computed on the output of each neuron. There are a number of different activation functions you could use. In practice, the three most widely used functions have been *sigmoid*, *tanh* and *ReLU*. Nonetheless, more complex activation functions exist (e.g. [2, 1]).



**Sigmoid** nonlinearity has form  $\sigma(x) = 1/(1 + e^{-x})$ . Sigmoid function squashes any real-valued input into range  $[0, 1]$ . It has seen frequent use historically, yet now it's rarely used. It has two major drawbacks:

- saturates and kill the gradient
- output is not zero centered

## Why using non-linear activations at all?

Composition of linear functions is a linear function. Without nonlinearities, neural networks would reduce to 1 layer logistic regression.

Let's say we have the following function ( $\phi$  represent nonlinear activations):

$$f(\mathbf{x}) = \phi(\mathbf{W}_2 \phi(\mathbf{W}_1 \mathbf{x}))$$

If we get rid of nonlinearities, this reduces to:

$$f(\mathbf{x}) = \mathbf{W}_2 \mathbf{W}_1 \mathbf{x} = \mathbf{W} \mathbf{x} \quad \text{where} \quad \mathbf{W} = \mathbf{W}_2 \mathbf{W}_1$$

which is clearly still linear.



**Softmax Classifier** generalizes Logistic Regression classifier to multi-class classification.

We first introduce the **softmax function**:

$$\text{softmax}_j(\mathbf{z}) = \frac{e^{z_j}}{\sum_k e^{z_k}}$$

It takes a vector of arbitrary real-valued scores  $\mathbf{z}$  and squashes it to a vector of values between zero and one that sum to one.

*e.g.*

$$\mathbf{z} = \begin{bmatrix} 1.2 \\ 5.1 \\ 2.7 \end{bmatrix} \quad \text{softmax}(\mathbf{z}) = \begin{bmatrix} 0.018 \\ 0.90 \\ 0.08 \end{bmatrix}$$

# Softmax Classifier ~D In output è incluso nelle loss. quindi non va fatto.

In the softmax classifier the scores of linear function mapping  $f(x_i, W) = Wx_i$  are interpreted as unnormalized log probabilities. To train the classifier we optimize the **cross-entropy loss**:


$$L = - \sum_i y_i^{true} \log(y_i^{pred})$$

where  $y_i^{true}$  is the ground truth output distribution and  $y_i^{pred}$  is the predicted output distribution. In practice,  $y_i^{true}$  is a *one-hot* vector selecting the output of the right label.

**Notice:** softmax classifier has the appealing property to produce an easy-to-interpret output, that is the normalized score confidence for each class.

The **goal** of this practice is to **implement a fully-connected neural network to perform 10-class classification** on the MNIST dataset.

**MNIST[3]** is a **database of handwritten digits** consisting of 60K training images and 10K testing images. All digits have been centered in 28x28 grayscale images.



## References

---

- [1] I. J. Goodfellow, D. Warde-Farley, M. Mirza, A. Courville, and Y. Bengio.  
**Maxout networks.**  
*arXiv preprint arXiv:1302.4389*, 2013.
- [2] K. He, X. Zhang, S. Ren, and J. Sun.  
**Delving deep into rectifiers: Surpassing human-level performance on imagenet classification.**  
In *Proceedings of the IEEE international conference on computer vision*, pages 1026–1034, 2015.

[3] Y. LeCun.

**The mnist database of handwritten digits.**

*<http://yann.lecun.com/exdb/mnist/>, 1998.*